

Tesla Stock Price Forecasting with Time Series Analysis

Let's Begin

Importing all the necessary libraries

```
In [289... # import important libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os
import sys
```

Importing dataset

```
In [290... data=pd.read_csv("TSLA.csv")
```

```
In [291... data.head()
```

```
Out[291...
   Date      Open      High      Low      Close      Volume  Dividends  S
0  2019-05-21  39.551998  41.480000  39.208000  41.015999   90019500         0
1  2019-05-22  39.820000  40.787998  38.355999  38.546001   93426000         0
2  2019-05-23  38.868000  39.894001  37.243999  39.098000  132735500         0
3  2019-05-24  39.966000  39.995998  37.750000  38.125999   70683000         0
4  2019-05-28  38.240002  39.000000  37.570000  37.740002   51564500         0
```

In this project, we will perform a Univariate Time Series Analysis

```
In [292... stock_data=data[["Date", "Close"]]
```

```
In [293... stock_data.head()
```

```
Out[293...]
      Date      Close
0  2019-05-21  41.015999
1  2019-05-22  38.546001
2  2019-05-23  39.098000
3  2019-05-24  38.125999
4  2019-05-28  37.740002
```

Checking information of the data

```
In [294...] stock_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 758 entries, 0 to 757
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  ---
0    Date    758 non-null      object
1   Close    758 non-null      float64
dtypes: float64(1), object(1)
memory usage: 12.0+ KB
```

Convert 'Date' feature into Date time data type using pandas

```
In [295...] stock_data["Date"] = pd.to_datetime(stock_data["Date"])
```

```
In [296...] stock_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 758 entries, 0 to 757
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  ---
0    Date    758 non-null      datetime64[ns]
1   Close    758 non-null      float64
dtypes: datetime64[ns](1), float64(1)
memory usage: 12.0 KB
```

Convert Date Column as Index Column

```
In [298...] stock_data = stock_data.set_index("Date")
```

```
In [353...] stock_data.head()
```

Out[353...

	Close
Date	
2019-05-21	41.015999
2019-05-22	38.546001
2019-05-23	39.098000
2019-05-24	38.125999
2019-05-28	37.740002

Exploratory Data Analysis(EDA)

Statistical Description


In [301... `stock_data.describe()`

Out[301...


	Close
count	758.000000
mean	485.531513
std	353.160353
min	35.793999
25%	112.323500
50%	488.125000
75%	762.142502
max	1229.910034

Visualizing Trends for Closing Price

```
In [355... # plotting close price
plt.style.use('ggplot')
plt.figure(figsize=(18,8))
plt.grid(True)
plt.xlabel('Dates', fontsize = 20)
plt.xticks(fontsize = 15)
plt.ylabel('Close Prices', fontsize = 20)
plt.yticks(fontsize = 15)
plt.plot(stock_data['Close'], linewidth = 3, color = 'blue')
plt.title('Tesla Stock Closing Price', fontsize = 20)
plt.show()
```



```
In [306... # Distribution of the close price
df_close = stock_data['Close']
df_close.plot(kind='kde',figsize = (18,8), linewidth= 3)
plt.xticks(fontsize = 15)
plt.grid("both")
plt.ylabel('Density', fontsize = 20)
plt.yticks(fontsize = 15)
plt.show()
```



Identifying if Time Series Data is Stationary or Non-Stationary by testing Visualization based test and Statistics based test

Using Visualization Test

```
In [307... # Moving Average
rolmean=stock_data["Close"].rolling(48).mean()
```

```
In [308... rolstd=stock_data["Close"].rolling(48).std()
```

```
In [ ]: plt.figure(figsize=(18, 8))

# Plot the closing prices
plt.plot(stock_data.Close, label='Closing Prices', color='blue')


# Plot the rolling mean
plt.plot(rolmean, label='Rolling Mean', color='orange')

# Plot the rolling standard deviation
plt.plot(rolstd, label='Rolling Std Dev', color='green')

# Add title and labels
plt.title('Stock Data with Rolling Statistics', fontsize=16)
plt.xlabel('Date', fontsize=14)
plt.ylabel('Price', fontsize=14)

# Add legend for clarity
plt.legend(loc='best', fontsize=12)

# Display the plot
plt.show()
```



ADF(Augmented Decay-Fuller Test)

```
In [310... #ADF(ada fullar test)
from statsmodels.tsa.stattools import adfuller
adft=adfuller(stock_data["Close"])
```

```
In [311... pd.Series(adft[0:4],index=["test stats","p-value","lag","data points"])
```


```
Out[311... test stats      -1.363009
p-value         0.599876
lag              9.000000
data points     748.000000
dtype: float64
```

- The most important value is P-Value

```
In [312... #Test for staionarity
def test_stationarity(timeseries):
    # Determing rolling statistics
    rolmean = timeseries.rolling(48).mean() # rolling mean
    rolstd = timeseries.rolling(48).std() # rolling standard deviation
    # Plot rolling statistics:
    plt.figure(figsize = (18,8))
    plt.grid('both')
    plt.plot(timeseries, color='blue',label='Original', linewidth = 3)
    plt.plot(rolmean, color='red', label='Rolling Mean',linewidth = 3)
    plt.plot(rolstd, color='black', label = 'Rolling Std',linewidth = 4)
    plt.legend(loc='best', fontsize = 20, shadow=True,facecolor='lightpink',
    plt.title('Rolling Mean and Standard Deviation', fontsize = 25)
    plt.xticks(fontsize = 15)
    plt.yticks(fontsize = 15)
    plt.show(block=False)

    print("Results of dickey fuller test")
    adft = adfuller(timeseries,autolag='AIC')
    # output for dft will give us without defining what the values are.
    # hence we manually write what values does it explains using a for loop
    output = pd.Series(adft[0:4],index=['Test Statistics','p-value','No. of
    for key,values in adft[4].items():
        output['critical value (%s)'%key] = values
    print(output)
```

```
In [313... test_stationarity(stock_data.Close)
```



```
Results of dickey fuller test
Test Statistics          -1.363009
p-value                 0.599876
No. of lags used        9.000000
Number of observations used 748.000000
critical value (1%)      -3.439123
critical value (5%)      -2.865412
critical value (10%)     -2.568832
dtype: float64
```

- From Visualization based test, we can see:
 1. Data is following Upward trend
 2. Moving Average is not constant
- From ADF test, we can see p value is greater than 0.5

Hence, Time Series data is Non-Stationary and we will accept the null hypothesis

Time Series Decomposition

```
In [314... from statsmodels.tsa.seasonal import seasonal_decompose
result=seasonal_decompose(stock_data[["Close"]],period=12)
```


```
In [315... result.seasonal
```

```
Out[315... Date
2019-05-21    -2.346452
2019-05-22     3.768884
2019-05-23    -0.777006
2019-05-24    -0.654226
2019-05-28    -2.737845
...
2022-05-16     2.149519
2022-05-17     1.323680
2022-05-18     1.837638
2022-05-19    -2.346452
2022-05-20     3.768884
Name: seasonal, Length: 758, dtype: float64
```

```
In [363... # Additive
from statsmodels.tsa.seasonal import seasonal_decompose
result=seasonal_decompose(stock_data[["Close"]],period=12,model="additive")

fig=plt.figure(figsize=(20,10))
fig=result.plot()
fig.set_size_inches(17,10)
```


<Figure size 2000x1000 with 0 Axes>



```
In [364... # Multiplicative
from statsmodels.tsa.seasonal import seasonal_decompose
result=seasonal_decompose(stock_data[["Close"]],period=12,model="multiplicat

fig=plt.figure(figsize=(20,10))
fig=result.plot()
fig.set_size_inches(17,10)
```


<Figure size 2000x1000 with 0 Axes>



Checking for outliers

```
In [319... # Checking for Outliers
import seaborn as sns
sns.boxplot(stock_data.Close)
```

```
Out[319... <Axes: ylabel='Close'>
```






- Data has no outliers

ACF (Autocorrelation Function) and PACF (Partial Autocorrelation Function) are statistical tools used in time series analysis to identify relationships between observations in a time series dataset.

```
In [365... from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
plot_acf(stock_data.Close)
plot_pacf(stock_data.Close)
```

Out[365...






Data Transformation

Converting Non-Stationary data into stationary data

```
In [322... df_close = df_close.diff()
df_close=df_close.dropna()
```

```
In [323... test_stationarity(df_close)
```



Kshitija Chilbule

```
Results of dickey fuller test
Test Statistics          -8.324564e+00
p-value                 3.498786e-13
No. of lags used        8.000000e+00
Number of observations used 7.480000e+02
critical value (1%)     -3.439123e+00
critical value (5%)     -2.865412e+00
critical value (10%)    -2.568832e+00
dtype: float64
```

- Now our data is stationary.
- Hence, will reject the null hypothesis

Model Building

In [326... *## Splitting data into training and testing dataset*

```
train_data = df_close[0:-60]
test_data  = df_close[-60:]
plt.figure(figsize=(18,8))
plt.grid(True)
plt.xlabel('Dates', fontsize=20)
plt.ylabel('Closing Prices', fontsize=20)
plt.xticks(fontsize=15)
plt.plot(train_data, "green", label='Train data', linewidth=5)
plt.plot(test_data, 'blue', label='Test data', linewidth=5)
plt.legend(fontsize=20, shadow=True, facecolor='lightpink', edgecolor='k')
```

Out[326... <matplotlib.legend.Legend at 0x1dbd2ec73b0>



Applying ARIMA Model on the Time Series

In [327... `import statsmodels.api as sm`
`from statsmodels.tsa.arima.model import ARIMA`
`from sklearn.metrics import mean_squared_error, mean_absolute_error`

```
In [328... history=[x for x in train_data]
```

```
In [329... model=ARIMA(history,order=(1,1,1))
```

```
In [330... model=model.fit()
```

```
In [331... model.summary()
```

```
Out[331... SARIMAX Results
```

Dep. Variable:	y	No. Observations:	698
Model:	ARIMA(1, 1, 1)	Log Likelihood	-3150.350
Date:	Sun, 15 Dec 2024	AIC	6306.700
Time:	19:26:27	BIC	6320.340
Sample:	0	HQIC	6311.974
	- 698		
Covariance Type:	opg		

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.2397	0.699	0.343	0.731	-1.129	1.609
ma.L1	-0.2713	0.690	-0.393	0.694	-1.623	1.080
sigma2	493.6767	11.690	42.231	0.000	470.765	516.588

Ljung-Box (L1) (Q):	0.00	Jarque-Bera (JB):	2111.30
Prob(Q):	0.98	Prob(JB):	0.00
Heteroskedasticity (H):	31.69	Skew:	0.09
Prob(H) (two-sided):	0.00	Kurtosis:	11.52

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [332... model.forecast()[0]
```

```
Out[332... 800.2043838663046
```

```
In [333... test_data[0]
```

```
Out[333... 809.8699951171875
```

```
In [334... mean_squared_error([test_data[0]],model.forecast())
```

Out[334...] 93.42404085319363

```
In [335...] np.sqrt(mean_squared_error([test_data[0]],model.forecast()))
```

Out[335...] 9.665611250882876

```
In [336...] def train_arima_model(X, y, arima_order):  
    # prepare training dataset  
    # make predictions list  
    history = [x for x in X]  
    predictions = list()  
    for t in range(len(y)):  
        model = ARIMA(history, order=arima_order)  
        model_fit = model.fit()  
        yhat = model_fit.forecast()[0]  
        predictions.append(yhat)  
        history.append(y[t])  
    # calculate out of sample error  
    rmse = np.sqrt(mean_squared_error(y, predictions))  
    return rmse
```

```
In [337...] def evaluate_models(data, test, p_values, d_values, q_values):  
    data = data.astype('float32')  
    best_score, best_cfg = float("inf"), None  
    for p in p_values:  
        for d in d_values:  
            for q in q_values:  
                order = (p,d,q)  
                try:  
                    rmse = train_arima_model(data, test, order)  
                    if rmse < best_score:  
                        best_score, best_cfg = rmse, order  
                    print('ARIMA%s RMSE=%.3f' % (order, rmse))  
                except:  
                    continue  
    print("Best ARIMA%s RMSE=%.3f" %(best_cfg, best_score))
```

```
In [338...] # evaluate parameters  
import warnings  
warnings.filterwarnings('ignore')  
p_values = range(0, 3)  
d_values = range(0, 3)  
q_values = range(0, 3)  
evaluate_models(train_data, test_data, p_values, d_values, q_values)
```

```

ARIMA(0, 0, 0) RMSE=457.414
ARIMA(0, 0, 1) RMSE=241.164
ARIMA(0, 0, 2) RMSE=161.913
ARIMA(0, 1, 0) RMSE=39.516
ARIMA(0, 1, 1) RMSE=39.482
ARIMA(0, 1, 2) RMSE=39.617
ARIMA(0, 2, 0) RMSE=57.835
ARIMA(0, 2, 1) RMSE=39.611
ARIMA(0, 2, 2) RMSE=39.580
ARIMA(1, 0, 0) RMSE=39.477
ARIMA(1, 0, 1) RMSE=39.449
ARIMA(1, 0, 2) RMSE=39.584
ARIMA(1, 1, 0) RMSE=39.475
ARIMA(1, 1, 1) RMSE=39.555
ARIMA(1, 1, 2) RMSE=39.935
ARIMA(1, 2, 0) RMSE=46.184
ARIMA(1, 2, 1) RMSE=39.573
ARIMA(1, 2, 2) RMSE=39.731
ARIMA(2, 0, 0) RMSE=39.440
ARIMA(2, 0, 1) RMSE=39.494
ARIMA(2, 0, 2) RMSE=39.598
ARIMA(2, 1, 0) RMSE=39.635
ARIMA(2, 1, 1) RMSE=39.759
ARIMA(2, 1, 2) RMSE=39.658
ARIMA(2, 2, 0) RMSE=45.781
ARIMA(2, 2, 1) RMSE=39.739
ARIMA(2, 2, 2) RMSE=39.733
Best ARIMA(2, 0, 0) RMSE=39.440

```

```

In [346... history = [x for x in train_data]
predictions = list()
conf_list = list()
for t in range(len(test_data)):
    model = ARIMA(history,order=(2,0,0))
    model_fit = model.fit()
    fc = model_fit.forecast(alpha = 0.05)
    predictions.append(fc)
    history.append(test_data[t])
print('RMSE of ARIMA Model:', np.sqrt(mean_squared_error(test_data, predicti


```

RMSE of ARIMA Model: 39.43995729937915

```


In [347... plt.figure(figsize=(18,8))
plt.grid(True)
plt.plot(range(len(test_data)),test_data, label = 'True Test Close Value', 1
plt.plot(range(len(predictions)), predictions, label = 'Predictions on test
plt.xticks(fontsize = 15)
plt.xticks(fontsize = 15)
plt.legend(fontsize = 20, shadow=True,facecolor='lightpink',edgecolor = 'k')
plt.show()

```



```
In [348... fc_series = pd.Series(predictions, index=test_data.index)
```

```
In [349... # Plot
plt.figure(figsize=(12,5), dpi=100)
plt.plot(train_data, label='Training', color = 'blue')
plt.plot(test_data, label='Test', color = 'green', linewidth = 3)
plt.plot(fc_series, label='Forecast', color = 'red')
plt.title('Forecast vs Actuals on test data')
plt.legend(loc='upper left', fontsize=8)
plt.show()
```



```
In [350... from statsmodels.graphics.tsaplots import plot_predict
fig = plt.figure(figsize=(18,8))
ax1 = fig.add_subplot(111)
plot_predict(result=model_fit,start=1, end=len(df_close)+60, ax = ax1)
plt.grid("both")
plt.legend(['Forecast', 'Close', '95% confidence interval'], fontsize = 20, sha
plt.show()
```





Applying SARIMA Model on the Time Series

```
In [351... # evaluate parameters
import warnings
warnings.filterwarnings('ignore')
history = [x for x in train_data]
predictions = list()
conf_list = list()
for t in range(len(test_data)):
    model = sm.tsa.statespace.SARIMAX(history, order = (0,1,0), seasonal_order=(0,1,0,0))
    model_fit = model.fit()
    fc = model_fit.forecast()
    predictions.append(fc)
    history.append(test_data[t])
print('RMSE of SARIMA Model:', np.sqrt(mean_squared_error(test_data, predictions)))
```

RMSE of SARIMA Model: 39.739481948001675

```
In [352... plt.figure(figsize=(18,8))
plt.title('Forecast vs Actual', fontsize = 25)
plt.plot(range(60), predictions, label = 'Predictions', linewidth = 4)
plt.plot(range(60), test_data, label = 'Close', linewidth = 4)
plt.legend(fontsize = 25, shadow=True, facecolor='lightpink', edgecolor = 'k')
```

Out[352... <matplotlib.legend.Legend at 0x1dbd2b738f0>



In []:

This notebook was converted with convert.ploomber.io